

16 fixed prompts | 512 output tokens | B=1 | TP=1 | CUDA Graph
Prefill + decode + offline API; one warmup excluded; no traces

